

Final Project

Alice Mee Seon Chung, Seulmin Yang

11/3/2017

Metropolis-Hastings

1.

Describing algorithm

1. Set a randomly selected start value from a uniform distribution (0,1)
2. Calculate ϕ_{new} from $Beta(c\phi_{old}, c(1 - \phi_{old}))$ using proporsal function.
3. Compute the acceptance probability. The Beta distribution is not symmetric, so we need to modify like below

$$Posterior = \frac{f(\phi_{new})}{f(\phi_i)}$$

$$Proposal = \frac{g(\phi_{new}|\phi_i)}{g(\phi_i|\phi_{new})}$$

$$acceptanceprobability = Posterior * Proposal$$

Final acceptance probability will be $\min(1, acceptanceprobability)$.

4. If the acceptance probability greater than a number from uniform distribution(0,1), then accept ϕ_{new} and assign $\phi_{i+1} = \phi_{new}$, otherwise, keep the current value and assign $\phi_{i+1} = \phi_i$.
5. Repeat 2~4 to finish all i times and get final chain.

R codes

```
# Set alpha and beta and startvalue
alpha <- 6
beta <- 4
startvalue <-runif(1)

# Proposal function in the problem to get new phi
proposalfunction <- function(c, phi_old){
  return(rbeta(1, c* phi_old, c*(1-phi_old)))
}

# Posterior function to compute the posetiror
posterior_beta<-function(x, alpha, beta){
  return (dbeta(x, alpha, beta))
}

# Proposal function to compute the proposal
proposal<-function(x, c, y){
  return (dbeta(x,c*y, c*(1-y)))
}
```

```

}

# Modify the code from Assignment2
run_betaMH <- function(startvalue, c, iterations){

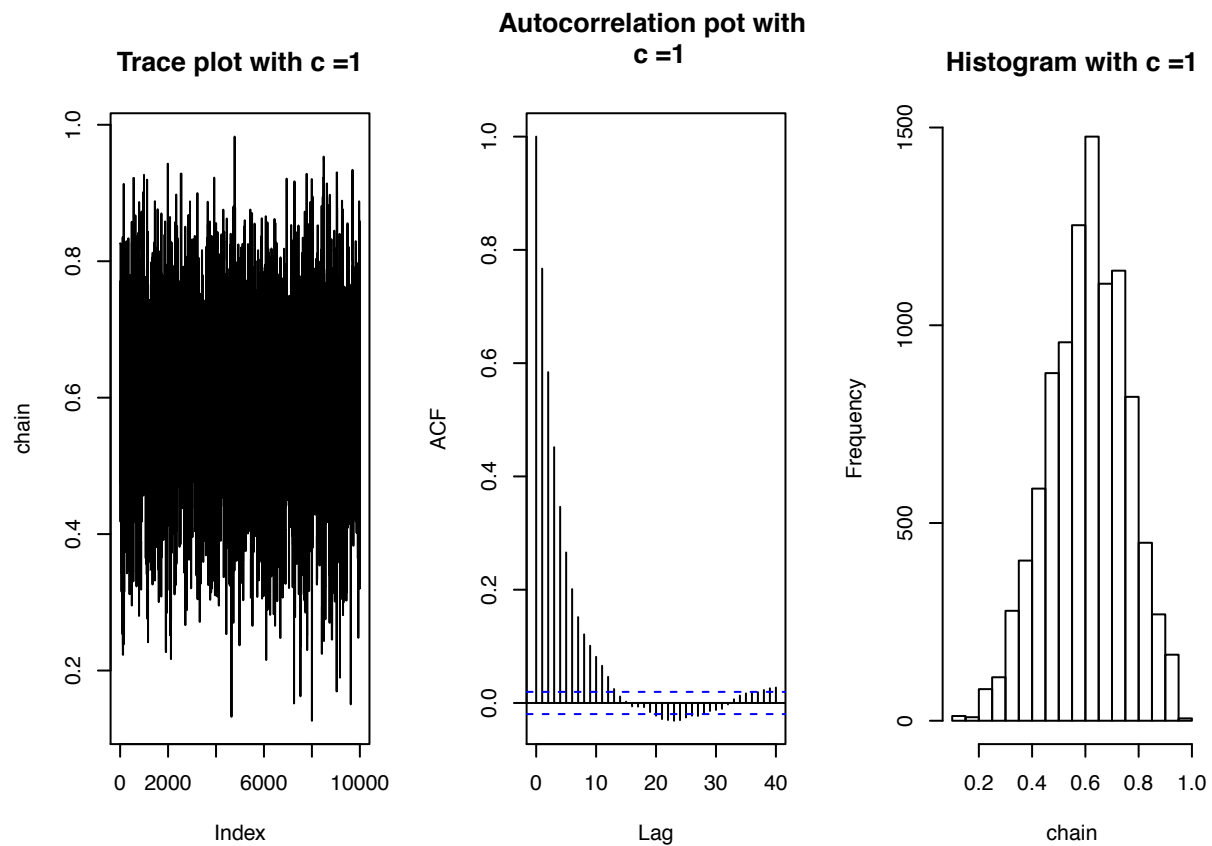
  chain = rep(0, iterations)
  chain[1] = startvalue
  for (i in 1:iterations){
    new_beta = proposalfunction(c, startvalue)
    posterior = posterior_beta(new_beta, alpha, beta) / posterior_beta(chain[i], alpha, beta)
    proposal = proposal(chain[i], c, new_beta) / proposal(new_beta, c, chain[i])
    probab = posterior * proposal
    # if acceptance probability is greater than the criteria, then accept new_beta
    # otherwise keep the current beta value
    if (runif(1) < min(1, probab)){
      chain[i+1] = new_beta
      startvalue = new_beta
    }else{
      chain[i+1] = chain[i]
    }
  }
  return(chain)
}

```

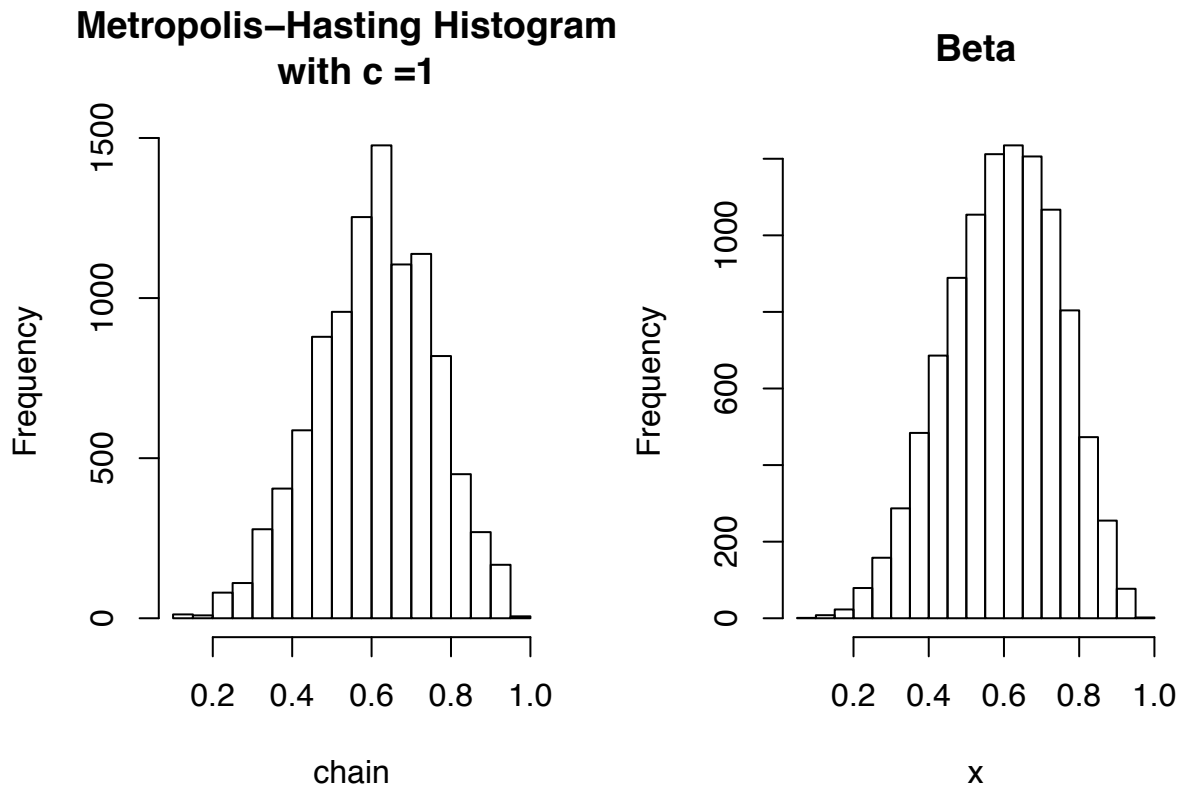
2.

```
alpha <- 6
beta <- 4
startvalue <- runif(1)
chain <- run_betaMH(startvalue, c=1, iterations = 10000)

par(mfrow=c(1,3))
plot(chain, type = 'l', main = 'Trace plot with c =1')
acf(chain, main = 'Autocorrelation pot with \n c =1')
hist(chain, main = 'Histogram with c =1')
```



```
x = rbeta(10000, 6, 4)
par(mfrow=c(1,2))
hist(chain, main = 'Metropolis-Hasting Histogram \n with c =1')
hist(x, main = 'Beta')
```



From above graphs, I think the performance of the sampler with c as 1 does not look fine. The traceplot showed the values of the parameter taken during the runtime of the chain. From two histograms plots, we can observe that the distribution of two histograms seems different, especially in the center part. It seems that the sampler result may not follow the $Beta(6, 4)$ distribution.

```
d<-ks.test(chain,x);d
```

```
## Warning in ks.test(chain, x): p-value will be approximate in the presence
## of ties

##
## Two-sample Kolmogorov-Smirnov test
##
## data: chain and x
## D = 0.037547, p-value = 1.508e-06
## alternative hypothesis: two-sided
```

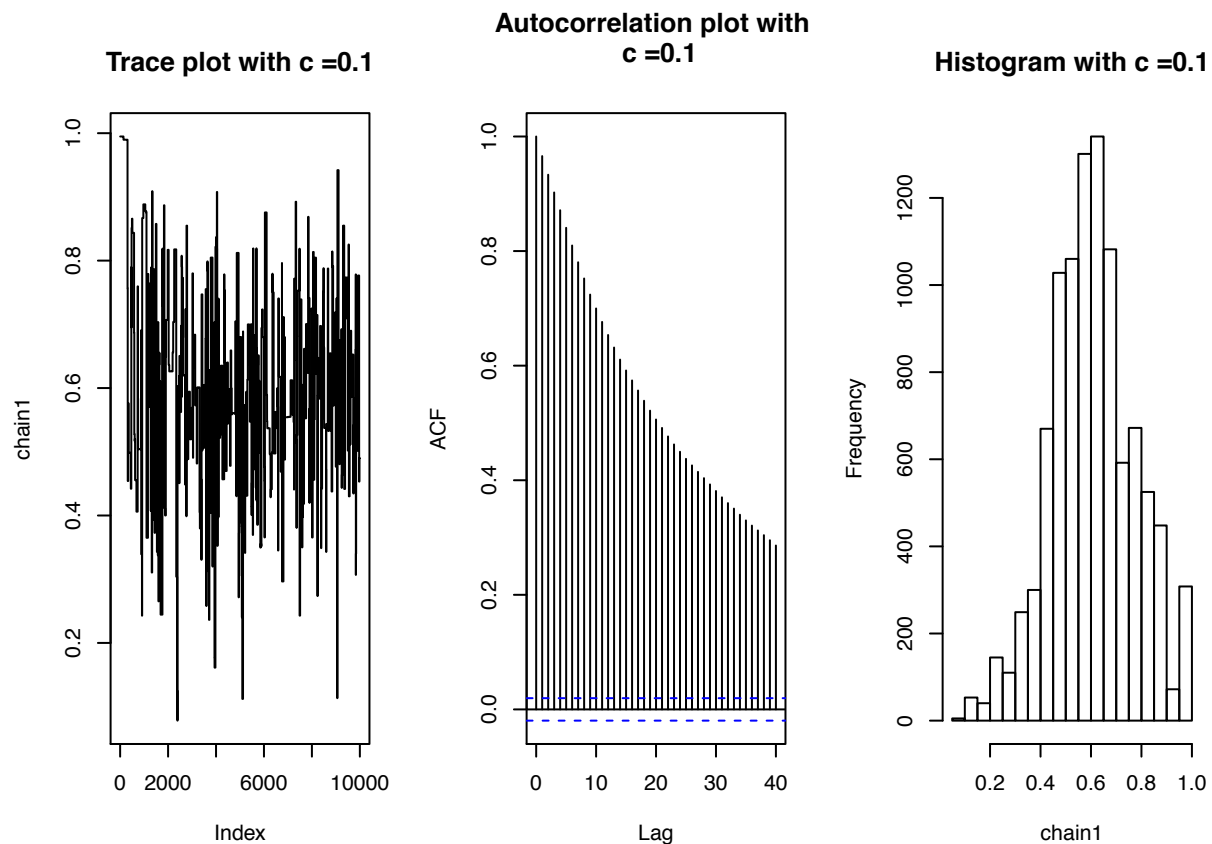
In addition to graphical tools, we also can check the result from Kolmogorov-Smirnov statistic. From the result, since p-value is less than $1.508e-06$, which is close to 0, we can reject the null hypothesis. Thus we can conclude that two distribution are not same.

Therefore the performance of the sampler with c as 1 is not good, and we can suspect that it may associated with the value of c .

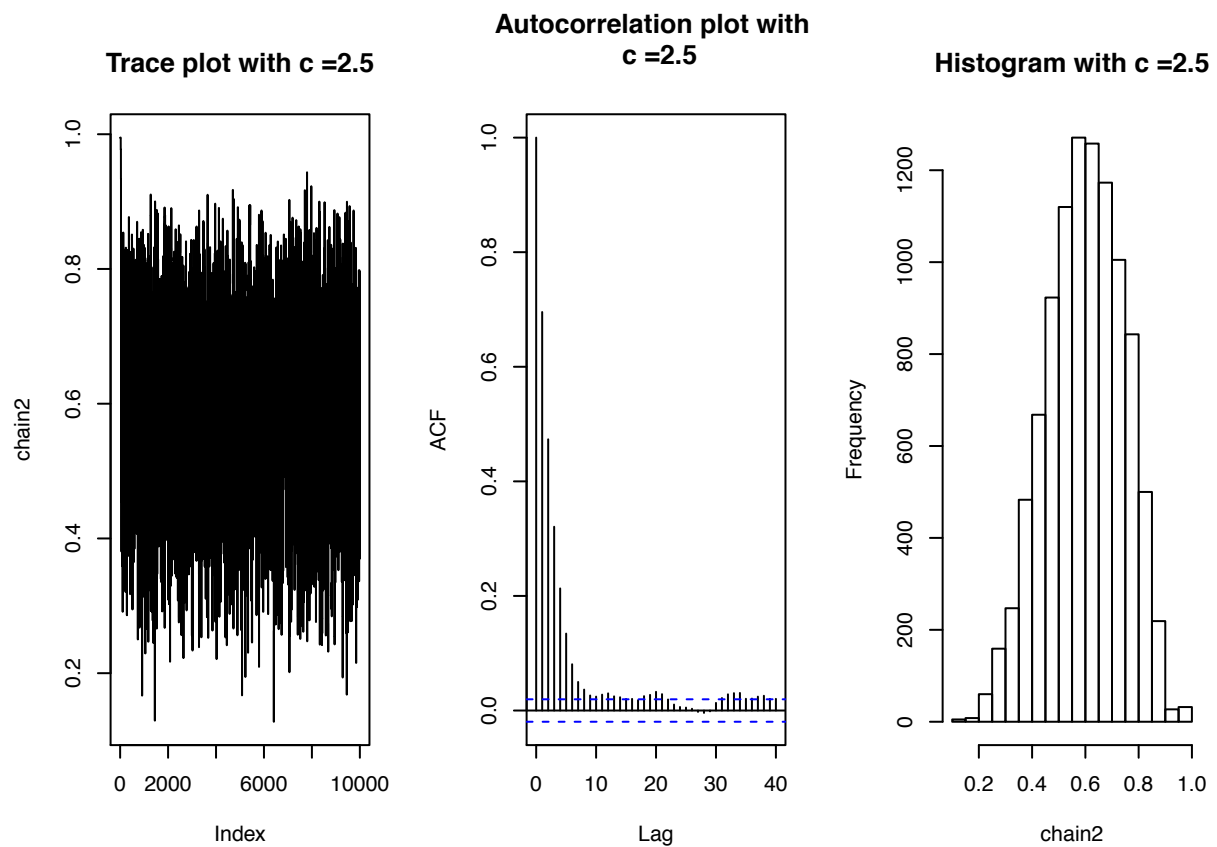
3.

```
alpha <- 6
beta <- 4
startvalue <- runif(1)
par(mfrow=c(1,3))

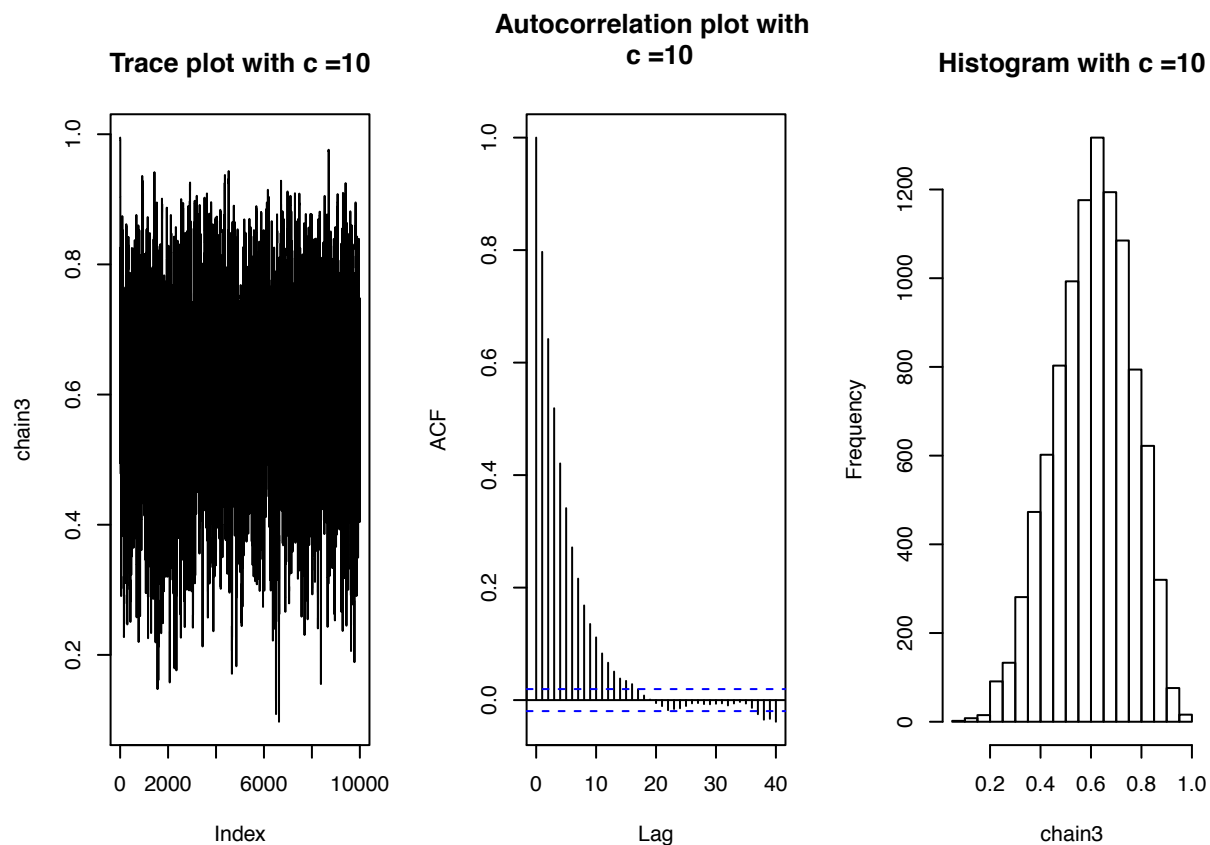
# c= 0.1
chain1<-run_betaMH(startvalue, c=0.1, iterations = 10000)
plot(chain1, type = 'l', main = 'Trace plot with c =0.1')
acf(chain1, main = 'Autocorrelation plot with \n c =0.1')
hist(chain1, main = 'Histogram with c =0.1')
```



```
# c = 2.5
chain2<-run_betaMH(startvalue, c= 2.5, iterations = 10000)
plot(chain2, type = 'l', main = 'Trace plot with c =2.5')
acf(chain2, main = 'Autocorrelation plot with \n c =2.5')
hist(chain2, main = 'Histogram with c =2.5')
```



```
# c = 10
chain3<-run_betaMH(startvalue, c=10, iterations = 10000)
plot(chain3, type = 'l', main = 'Trace plot with c =10')
acf(chain3, main = 'Autocorrelation plot with \n c =10')
hist(chain3, main = 'Histogram with c =10')
```



We rerun the sampler with $c = 0.1$, $c = 2.5$, and $c = 10$ and plot s , autocorrelation plots and histograms for these samplers. If we see these graphs all together, the one thing that we can see obviously is the differences in autocorrelation plots. When $c = 0.1$, the autocorrelation plot shows the most ACF compared with $c = 2.5$ and $c = 10$. From the histograms, we can see that the histogram with $c = 0.1$ looks different from the histograms with the histograms of $c = 2.5$ and $c = 10$ and the histograms with $c = 2.5$ and $c = 10$ seems similar to the target distribution. From autocorrelation plots, the the graph with $c = 2.5$ is less lagging and has less ACF than $c = 10$, and the histograms seems very similar to the target distribution. Thus I think the sampler with $c = 2.5$ is most effective at drawing from the target distribution.

In addition to graphical tools, we also can check the result from Kolmogorov-Smirnov statistic for these three samplers.

```
# c = 0.1
ks.test(chain1,x)
```

```
## Warning in ks.test(chain1, x): p-value will be approximate in the presence
## of ties
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: chain1 and x
## D = 0.059886, p-value = 5.551e-16
## alternative hypothesis: two-sided
```

```
# c = 2.5
ks.test(chain2,x)
```

```
## Warning in ks.test(chain2, x): p-value will be approximate in the presence
## of ties
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: chain2 and x
## D = 0.015034, p-value = 0.2084
## alternative hypothesis: two-sided
```

```
# c = 10
ks.test(chain3,x)
```

```
## Warning in ks.test(chain3, x): p-value will be approximate in the presence
## of ties
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: chain3 and x
## D = 0.032842, p-value = 4.136e-05
## alternative hypothesis: two-sided
```

From the results, since only the p-value of sampler with $c = 2.5$ is greater than 0.05, we can not reject the null hypothesis. So we can conclude that the sampler with $c = 2.5$ follow the target distribution, $Beta(6, 4)$.

As c increases, we can get the distribution that follows closer to the target distribution efficiently. However, it is not always true that large c is optimal option because it can increase lags. To increase efficiency and to optimize the process, we could include the appropriate numbers of draws needed for thinning and burn-in instead of simply increasing c .

Gibbs sampling

1.

Algorithms

From the marginal distribution, we can find the density function, cumulative density function and inverse cumulative density function. Here, B is a known positive constant.

Marginal ditribution

$$p(x | y) \propto ye^{-yx}, 0 < x < B$$

$$p(y | x) \propto xe^{-yx}, 0 < y < B$$

Density function

$$f(x|y) = \frac{ye^{-yx}}{1 - e^{-By}}, 0 < x < B$$

$$f(y|x) = \frac{xe^{-yx}}{1 - e^{-Bx}}, 0 < y < B$$

Cumulative density function

$$F(x|y) = \frac{1 - e^{-yx}}{1 - e^{-By}}, 0 < x < B$$

$$F(y|x) = \frac{1 - e^{-yx}}{1 - e^{-Bx}}, 0 < y < B$$

Inverse of Cumulative density function

$$F^{-1}(p_1|y) = \left(\frac{-1}{y}\right) * \log(1 - p_1(1 - e^{-By})), 0 \leq p_1 \leq 1$$

$$F^{-1}(p_1|x) = \left(\frac{-1}{x}\right) * \log(1 - p_1(1 - e^{-Bx})), 0 \leq p_1 \leq 1$$

We use inverse transform sampling to generate the sample from the conditional $p(x|y)$, so we use random number from uniform(0,1) to get x using **Inverse of Cumulative density function** as we define above. Same method will apply to the case of y .

1. Set first x with a random number from uniform(0,1) and first y with a random number from uniform(0,1).
2. Generate new x using inverse function with given y and new y using inverse function with given x .
3. In the next step, new x and new y in the previous step will be the given x and the given y to generate new x and new y .
4. Repeat 2-3 to finish all T times and get final matrix with the sample.

Here, we set thin value as 10 (default).

NOTE

Python codes and the solutions for later problems are written in jupyter notebook with pdf form. Those pages will attach from next page.

Gibbs Sampling

November 3, 2017

1 Python codes

```
In [1]: import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline

def pxy(a, b):
    '''
    Marginal distribution function.
    This function can be used into both  $p(x|y)$  and  $p(y|x)$ .
    '''
    return a * np.exp((-a)*b)

def densityfxy(a, b, B):
    '''
    Density function in an interval  $(0,B)$ .
    This function can be used into both  $f(x|y)$  and  $f(y|x)$ .
    '''
    return (b*np.exp((-a)*b))/(1-np.exp((-B)*b))

def cdffxy(a, b, B):
    '''
    Cumulative density function in an interval  $(0,B)$ .
    This function can be used into both  $F(x|y)$  and  $F(y|x)$ .
    '''
    return (1-np.exp((-a)*b))/(1-np.exp((-B)*b))

def inversefxy(a ,B):
    '''
    Inverse function in an interval  $(0,B)$ .
    This function can be used into both  $F^{-1}(p1|y)$  and  $F^{-1}(p1|x)$ ,
    where  $0 < p1 < 1$ .
    Let  $p1$  be the r.v from uniform distribution.
    '''
    p1 = random.uniform(0,1)
    return (-1/a)*(np.log(1-p1*(1-np.exp((-B)*a))))
```

```

In [2]: # Gibbs sampling method
def gibbs(N, B, thin):
    matrix = np.zeros((N, 3))
    matrix[:,0] = np.arange(1,N+1)
    # select random numbers from uniform distribution
    x = random.uniform(0,1)
    y = random.uniform(0,1)
    for i in range(0, N):
        for j in range(0, thin):
            # generate updated x value
            x = inversefxy(y,B)
            # generate updated y value
            y = inversefxy(x,B)
            # assign updated x,y values to column 2,3
            matrix[i,1:3] = [x,y]
    return matrix

```

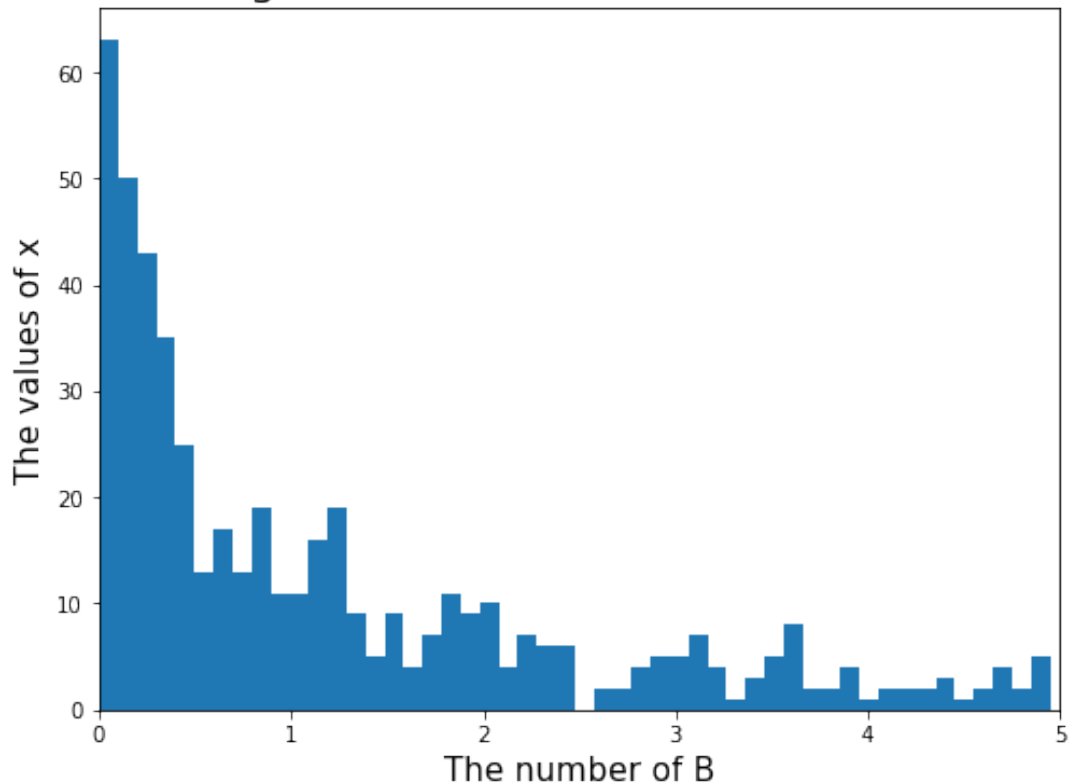
2 2

```

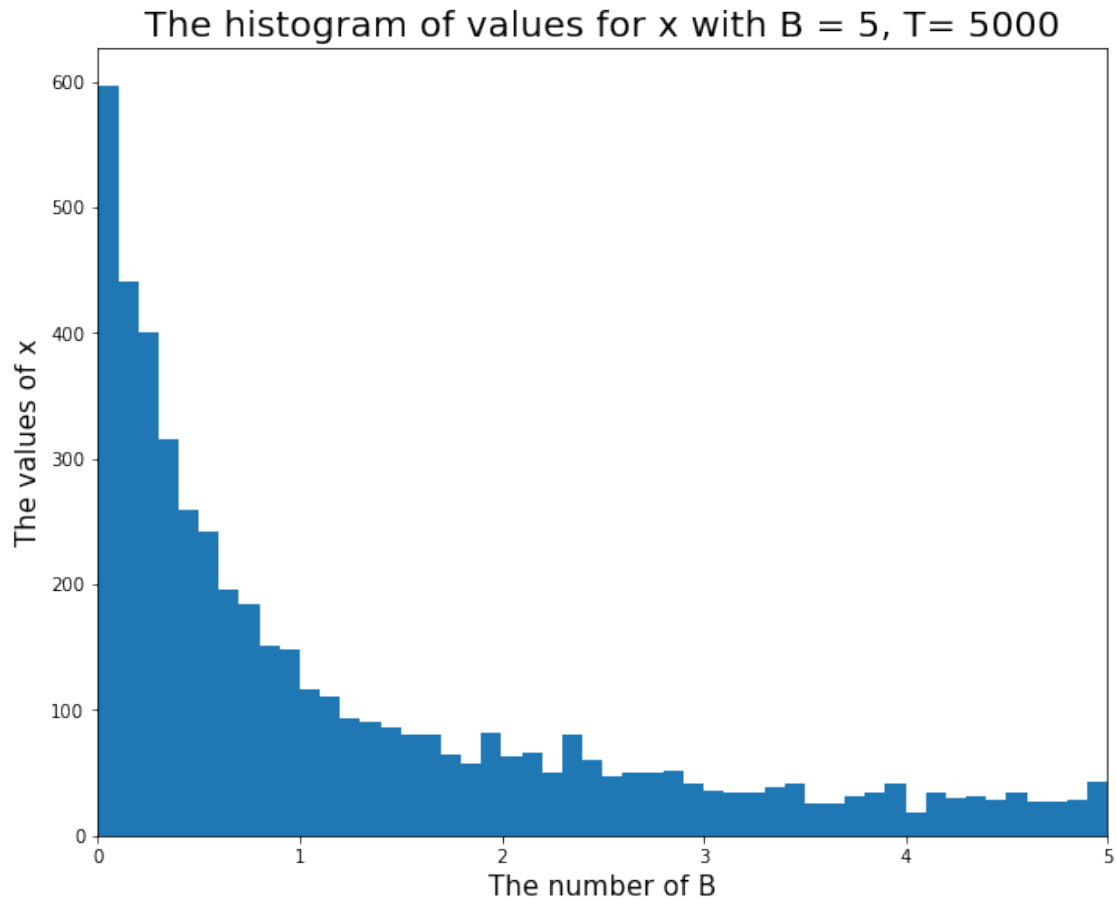
In [3]: # Plot histogram with T = 500
N, B, thin = 500, 5, 10
result1 = gibbs(N, B, thin)
plt.subplots(figsize=(8, 6))
plt.hist(result1[:,1], bins = 50)
plt.title ('The histogram of values for x with B = 5, T= 500',
           fontsize= 20)
plt.xlabel('The number of B', fontsize= 15)
plt.ylabel('The values of x', fontsize= 15)
plt.xlim(0,5)
plt.show()

```

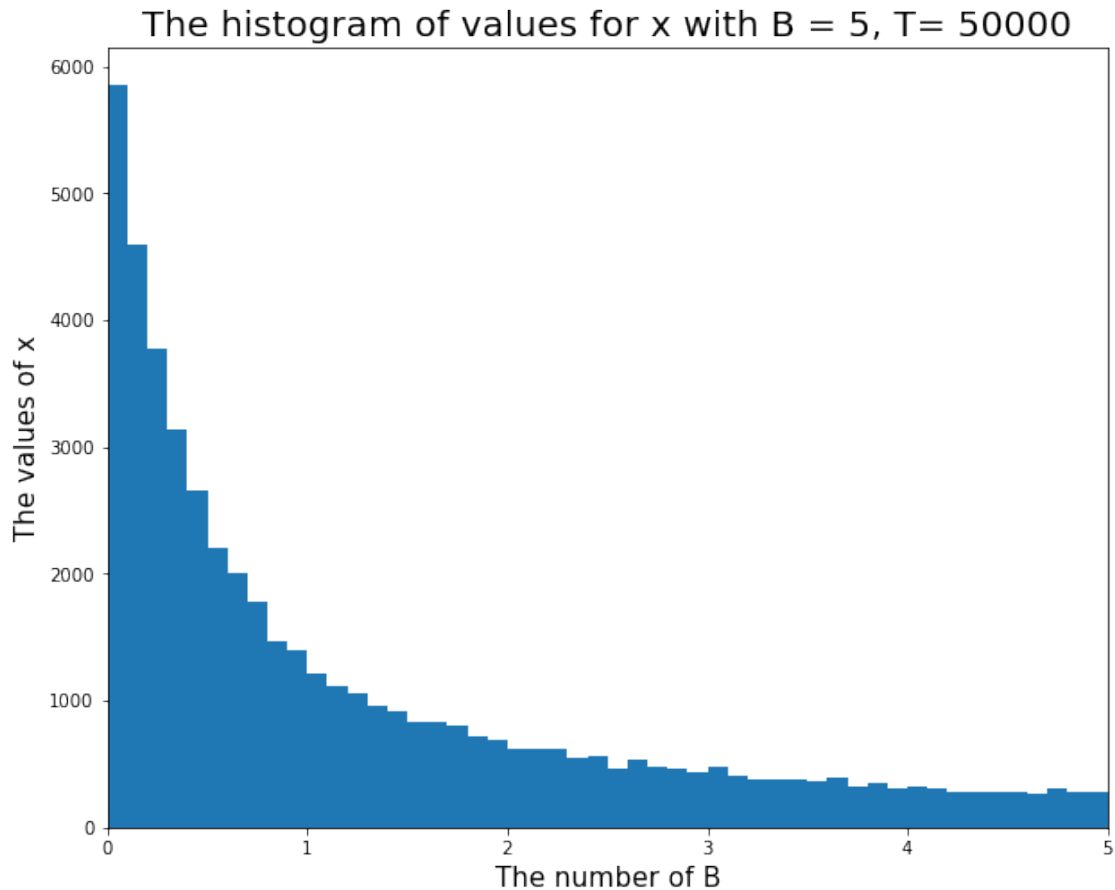
The histogram of values for x with $B = 5$, $T = 500$



```
In [4]: # Plot histogram with T = 5000
N, B, thin = 5000, 5, 10
result2 = gibbs(N, B, thin)
plt.subplots(figsize=(10, 8))
plt.hist(result2[:,1],bins = 50)
plt.title ('The histogram of values for x with B = 5, T= 5000',
           fontsize= 20)
plt.xlabel('The number of B',fontsize= 15)
plt.ylabel('The values of x',fontsize= 15)
plt.xlim(0,5)
plt.show()
```



```
In [5]: # Plot histogram with T = 50000
N, B, thin = 50000, 5, 10
result3 = gibbs(N, B, thin)
plt.subplots(figsize=(10, 8))
plt.hist(result3[:,1],bins = 50)
plt.title ('The histogram of values for x with B = 5, T= 50000',
          fontsize= 20)
plt.xlabel('The number of B',fontsize= 15)
plt.ylabel('The values of x',fontsize= 15)
plt.xlim(0,5)
plt.show()
```



From these three histograms, we can see that as T increases, the histograms become smoother.

3 3.

```
In [18]: random.seed(200)
         result500 = gibbs(500,5,10)
         result5000 = gibbs(5000,5,10)
         result50000 = gibbs(50000,5,10)

         # calculate the means from each iterations.
         estimated_expectation_500 = np.mean(result500[:,1])
         estimated_expectation_5000 = np.mean(result5000[:,1])
         estimated_expectation_50000 = np.mean(result50000[:,1])
```

```
In [19]: estimated_expectation_500
```

```
Out[19]: 1.2776584241390458
```

```
In [20]: estimated_expectation_5000
```

```
Out [20]: 1.2662195687732802
```

```
In [21]: estimated_expectation_50000
```

```
Out [21]: 1.2584644918562236
```

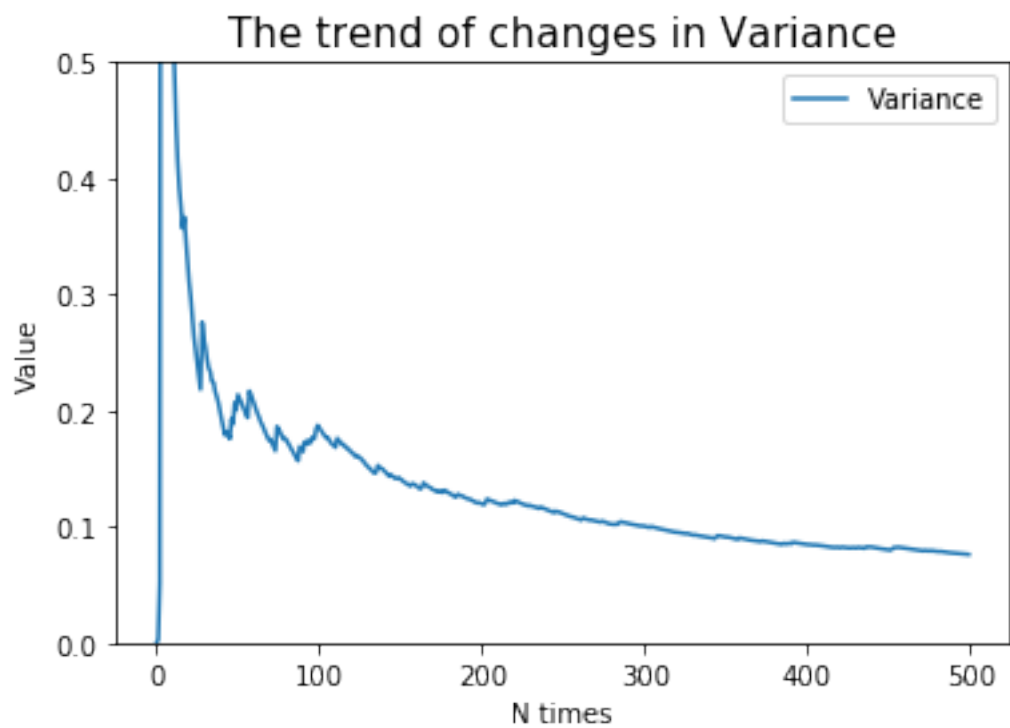
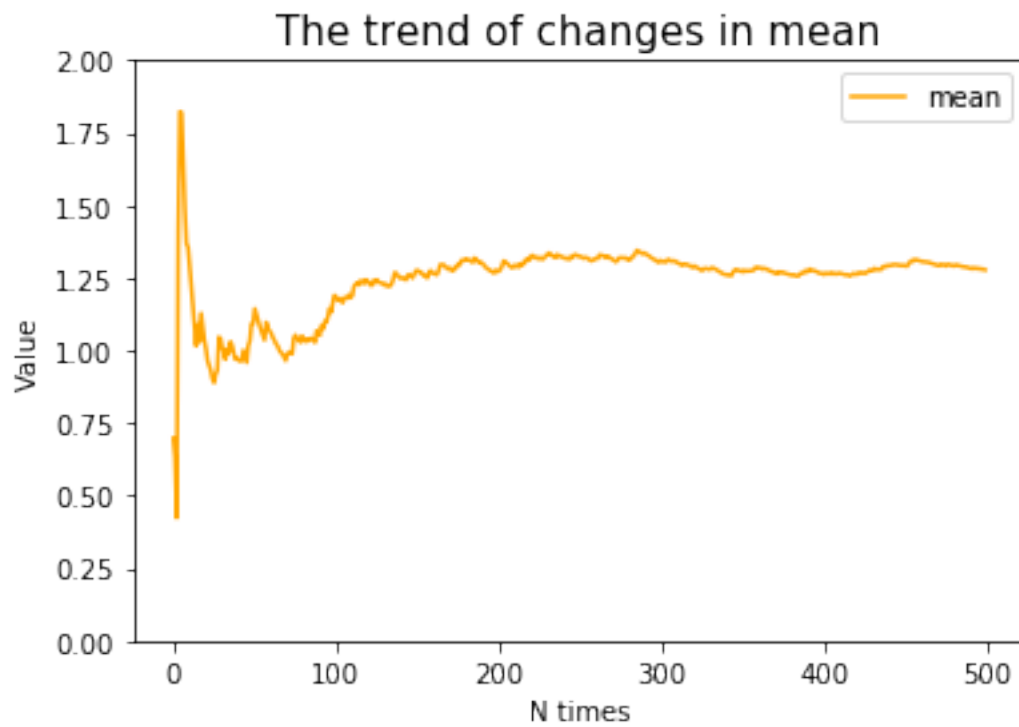
As we can see from the results, the expectation changes 1.277658 to 1.266219 to 1.258464 by using 500, 5000 and 50000 samples respectively. The expectation become stable as T increases like 500, 5000 and 50000 samples from the sampler. We can see this trend by using graphs.

```
In [25]: # define function do simulation and make graph at the same time
# the function takes T times as argument
def simulation_and_graph(N):
    simulation_matrix = np.zeros((N,3))
    simulation_matrix[:,0] = np.arange(N)
    simulation = gibbs(N, 5, 10)
    for i in range(N):
        mean = np.mean(simulation[0:i+1,1])
        var = np.var(simulation[0:i+1,1])/np.sqrt(i+1)
        simulation_matrix[i,1:3] = [mean, var]

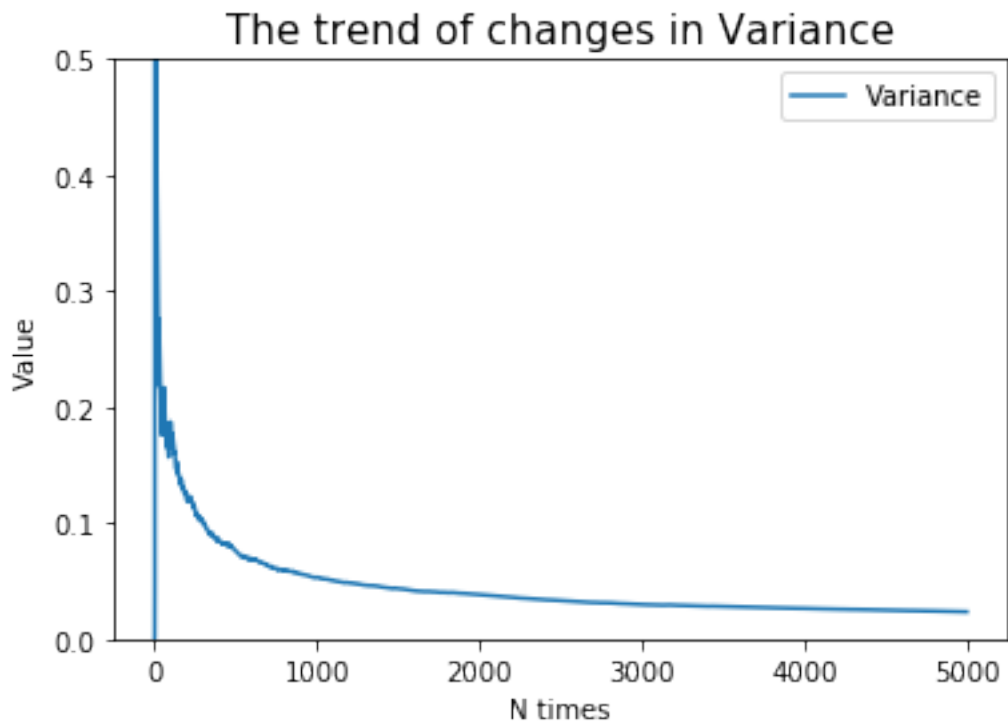
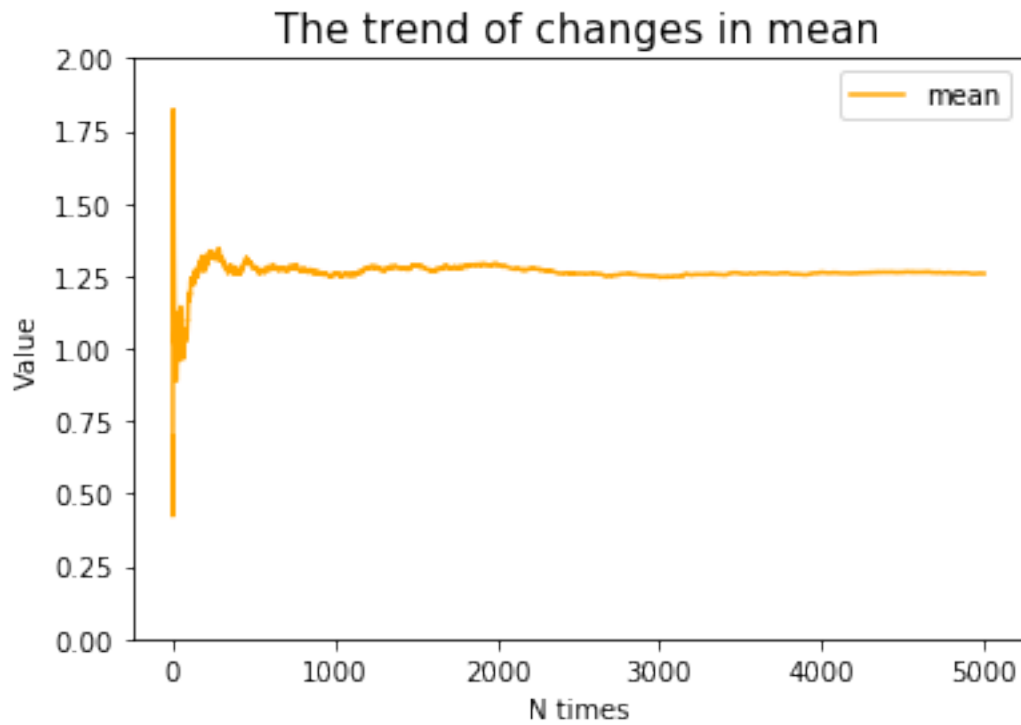
    plt.subplots(figsize=(6, 4))
    plt.plot(simulation_matrix[:,1], label = 'mean', color = 'orange')
    plt.title ('The trend of changes in mean',fontsize= 15)
    plt.xlabel('N times',fontsize= 10)
    plt.ylabel('Value',fontsize= 10)
    plt.legend(loc='upper right')
    plt.ylim(0,2)
    plt.legend

    plt.subplots(figsize=(6, 4))
    plt.plot(simulation_matrix[:,2], label = 'Variance')
    plt.title ('The trend of changes in Variance',fontsize= 15)
    plt.xlabel('N times',fontsize= 10)
    plt.ylabel('Value',fontsize= 10)
    plt.legend(loc='upper right')
    plt.ylim(0,0.5)
    plt.legend

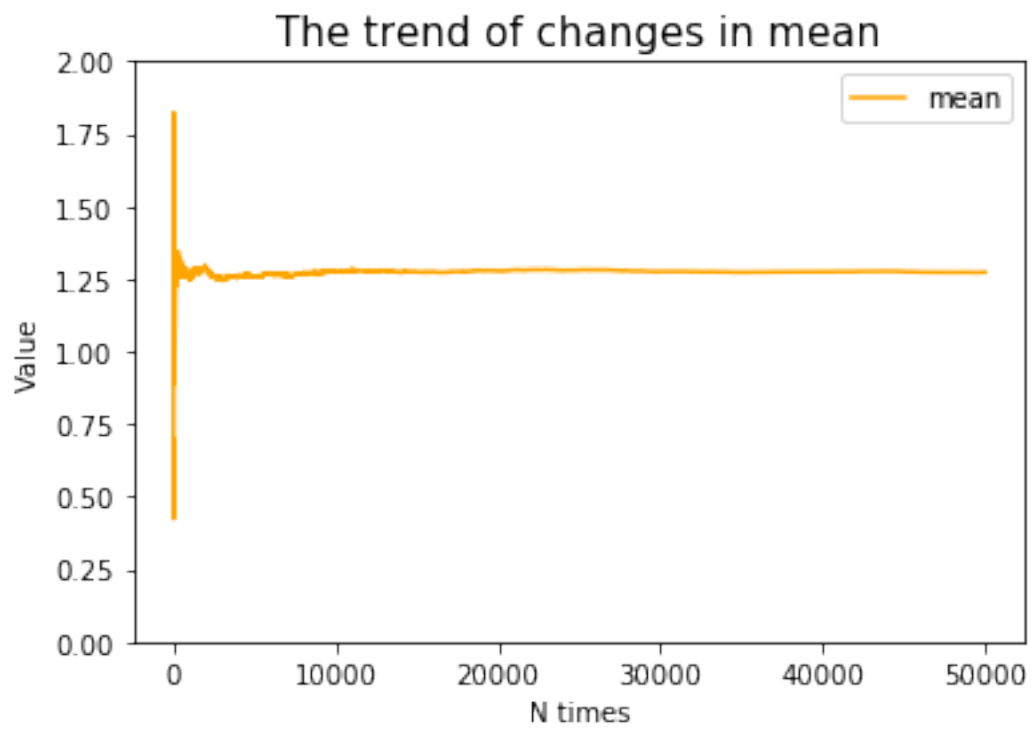
In [29]: # T = 500
random.seed(200)
simulation_and_graph(500)
```

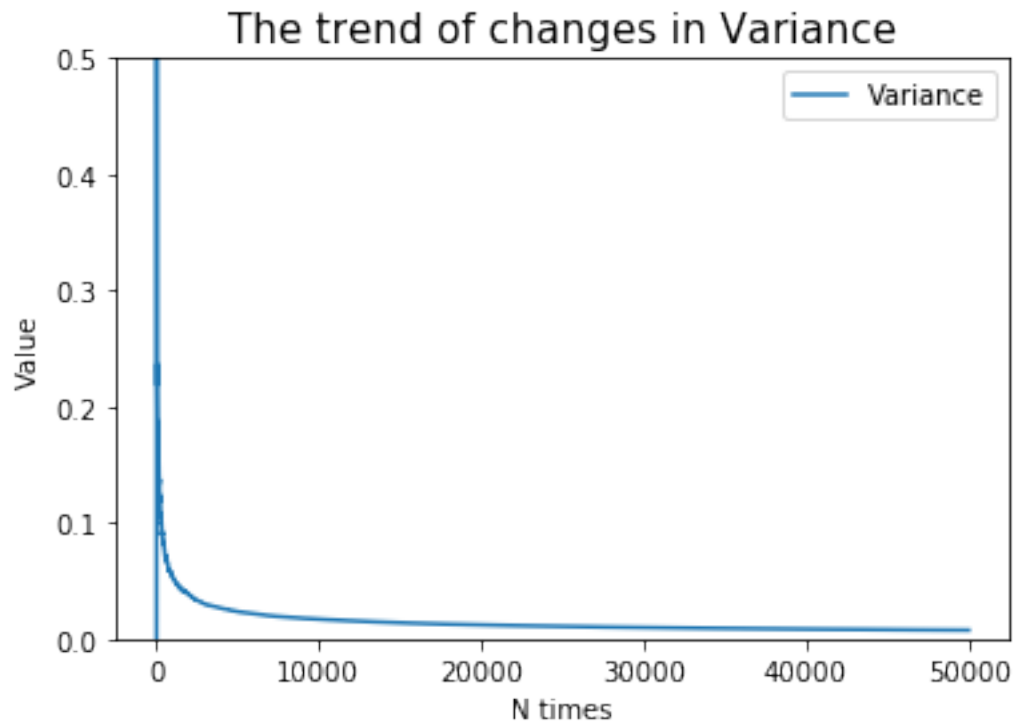



```
In [30]: # T = 5000
random.seed(200)
simulation_and_graph(5000)
```



```
In [31]: #  $T = 500$   
random.seed(200)  
simulation_and_graph(50000)
```





As you see in the above graphs that shows how means and variance change as T increases, the trend of mean and variance varies a lot at first and becomes stable and converges. Especially variance becomes small as T increases.

In []:

K-means

1.

Algorithms

1. Select the number of clusters and randomly select the first random centers from the data.
2. For every points in the data, calculate the euclidean distance with the cluster centers, choose the minimum distance and assign into the according group.
3. Recalculate the means of cluster centers within each groups and set as new cluster centers.
4. Compare the previous cluster centers and new cluster centers, if the centers changed then repeat until the centers do not change.
5. Get and return the final cluster result.

R codes

```
data(wine, package = 'rattle.data')
data.wine <- wine[-1]
data.iris <- iris[-5]

# the number of cluster groups
cluster_num = 3

mykmeans<-function(data, cluster_num){
  data <- as.matrix(data)
  rows <- nrow(data)
  cols <- ncol(data)
  center <- matrix(0, cluster_num, cols)
  resultmatrix <- matrix(0, rows, 2)
  euclidist <- matrix(0, 1, cluster_num)
  # randomly select the first sample
  random <- sample(rows, cluster_num, replace = FALSE)
  ## STEP 1
  # update center matrix
  for (i in 1:cluster_num){
    center[i,] <- data[random[i,]] # use random to get first center points
  }
  previouscenter <- matrix(0, cluster_num, cols)
  # within this while loop, compare current centers and previous centers.
  # if those points are changed, we do STEP2 and STEP3 again.
  # if those points are not changes, then we return the final results.
  while(sum(abs(previouscenter)) != sum(abs(center))){
    ## STEP 2
    # update euclidian distance matrix for every rows
    for (i in 1:rows){
      for (j in 1:cluster_num){
        # calculate euclidean distance measure
        euclidist[,j] <- sqrt(sum((data[i,]-center[j,])^2))
      }
      # assign every rows to one of cluster groups and record every calculated distance
      resultmatrix[i,] = c(min <- which.min(euclidist),
        euclidist[min <- which.min(euclidist)])
    }
  }
}
```

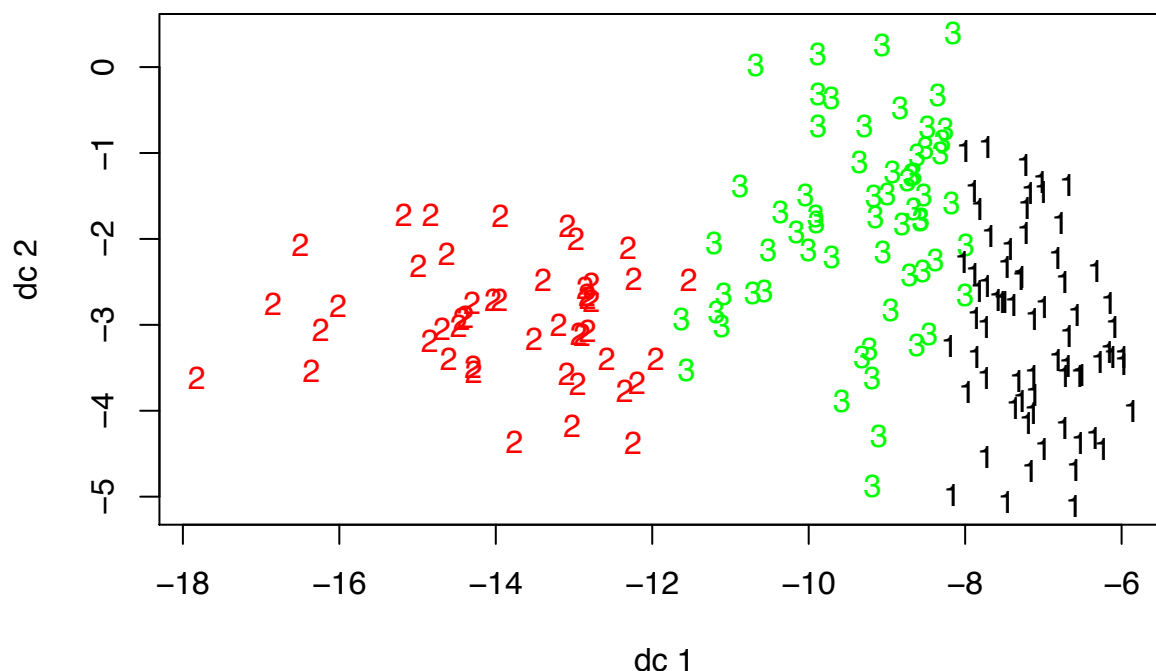
```

}
## STEP 3
previouscenter <- center
# calculate a new center for every cluster group
for (i in 1:cluster_num){
  index = which(resultmatrix[,1]==i, arr.in=TRUE)
  center[i,] = colMeans(data[index,])
}
}
return(resultmatrix)
}

```

Wine data

```
kmeans_wine <- mykmeans(data.wine, 3)
plotcluster(data.wine, as.list(kmeans_wine[,1]))
```



As we can see in the above cluster plot, it is separated into three groups but the boundary of each cluster is quite blurry. Some data points overlaps each other. So the clusters here doesn't seem to be well-separated.

Let's develop a method to quantify how well my algorithm's clusters correspond to the three wine types. I simply want to compare with the original dataset. In the dataset, there is **Type** variable and it is also clustered into three groups. So we can compare the cluster result with original data, find and count the matched pairs and divide by the number of whole data points. It is simple ratio and not a complex method to compute.

```
correct_ratio = sum(kmeans_wine[,1] == wine[1])/length(kmeans_wine[,1])
```

So if we calculate with above cluster results, then this correct ratio will be 0.1685393. It is less than 20%, so we can not say that it is good clustering and it explains why we have overlaps in cluster plot.

Second method is using built-in R method **randIndex**. This method computes the Rand index, a measure of the similarity between two clusterings.

```
table <- table(wine$Type, kmeans_wine[,1])
ari<- randIndex(table); ari
```

```
##          ARI
## 0.3711137
```

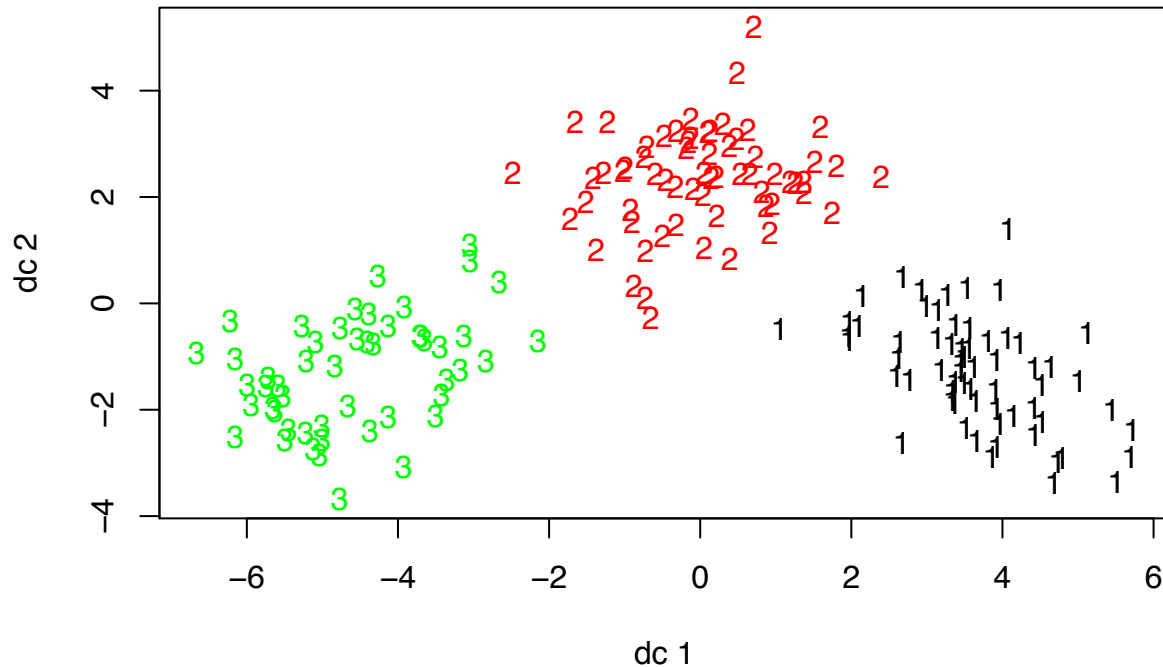
The results is 0.3711137 and it is not high value to say two clusterings are similar. Thus, my kmeans algorithm does not separate clusters well.

```
#scaled
scaled.wine = scale(wine[-1])
```

Now we will use scale function to scale data. This function will calculate the mean and standard deviation of the entire vector, then scale each element by those values by subtracting the mean and dividing by the sd. So

it minimize the effect of the outliers in the dataset. If we do not scale this data observations, then outliers or unusual data observations will affect and distort the results from whatever algorithms we use to analyze the data. Thus if we use scale function here, our clustering results will be more precise and have less errors.

```
kmeans_wine.scaled <- mykmeans(scaled.wine, 3)
plotcluster(scaled.wine, as.list(kmeans_wine.scaled[,1]))
```



This time we used scaled wine data and cluster plot seems work well. There is no overlapping at the boundaries of cluster groups. So the clusters here seem to be well-seperated into three groups.

We can calculate correct ratio.

```
correct_ratio1 = sum(kmeans_wine.scaled[,1] == wine[1])/length(kmeans_wine.scaled[,1])
```

So if we calculate with above cluster results, then this correct ratio will be 0.9550562. It is very close to 1, so we can say that it is good clustering. Also compare to previous result with unscaled data, it increases almost two times.

We also can use second method randIndex.

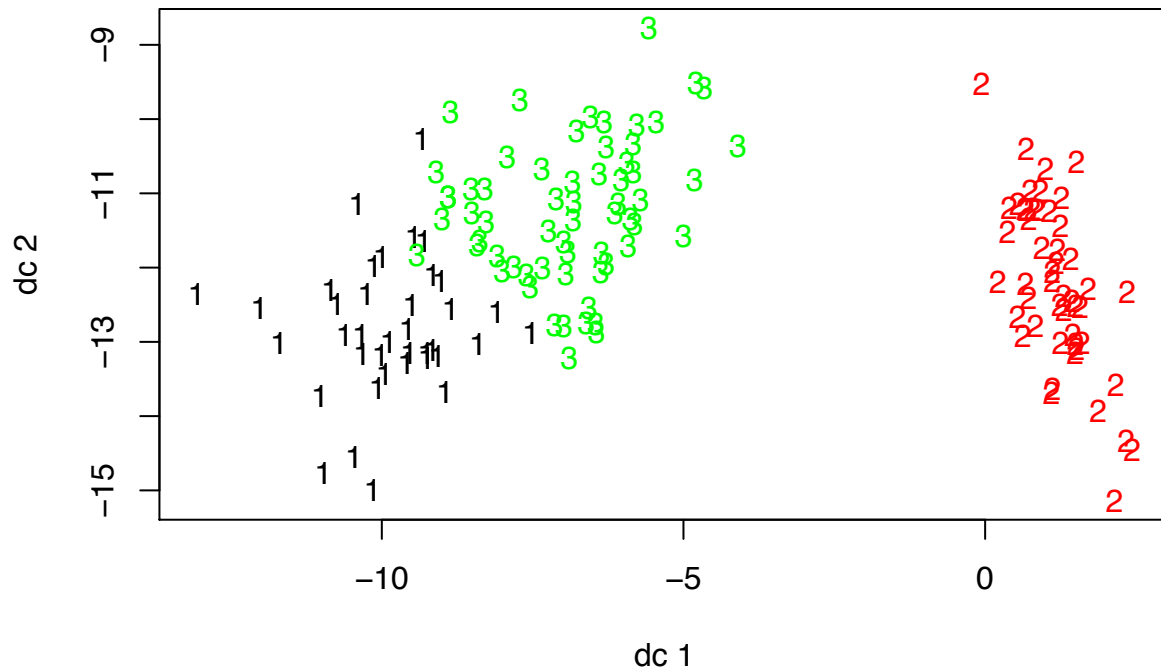
```
table1 <- table(wine$Type, kmeans_wine.scaled[,1])
ari1<- randIndex(table1); ari1
```

```
##          ARI
## 0.8635988
```

The results is 0.8635988 and it is high value to say two clustertings are similar. Thus, my kmeans algoritgm seperates clusters well with scaled wine data.

Iris data

```
kmeans_iris <- mykmeans(data.iris, 3)
plotcluster(data.iris, as.list(kmeans_iris[,1]))
```



As we can see in above cluster plot with unscaled iris data, it is separated into three groups but the boundary of cluster 3 and 2 are vague. Here also, Some data points overlaps each other. So the clusters here doesn't seem to be well-separated.

In the iris dataset, there is **Species** variable and it is also clustered in three groups. So we can use **simple ratio** that we define earlier.

```
iris$Species<-as.numeric(iris$Species, levels = sort(iris$Species))
correct_ratio3 = sum(kmeans_iris[,1] == iris[5])/length(kmeans_iris[,1])
```

So if we calculate with above cluster results, then this correct ratio will be 0.0933333. It is less than 10%, so we can not say that it is good clustering and it explains why we have overlaps in cluster plot.

We can use second method, built-in R method **randIndex**.

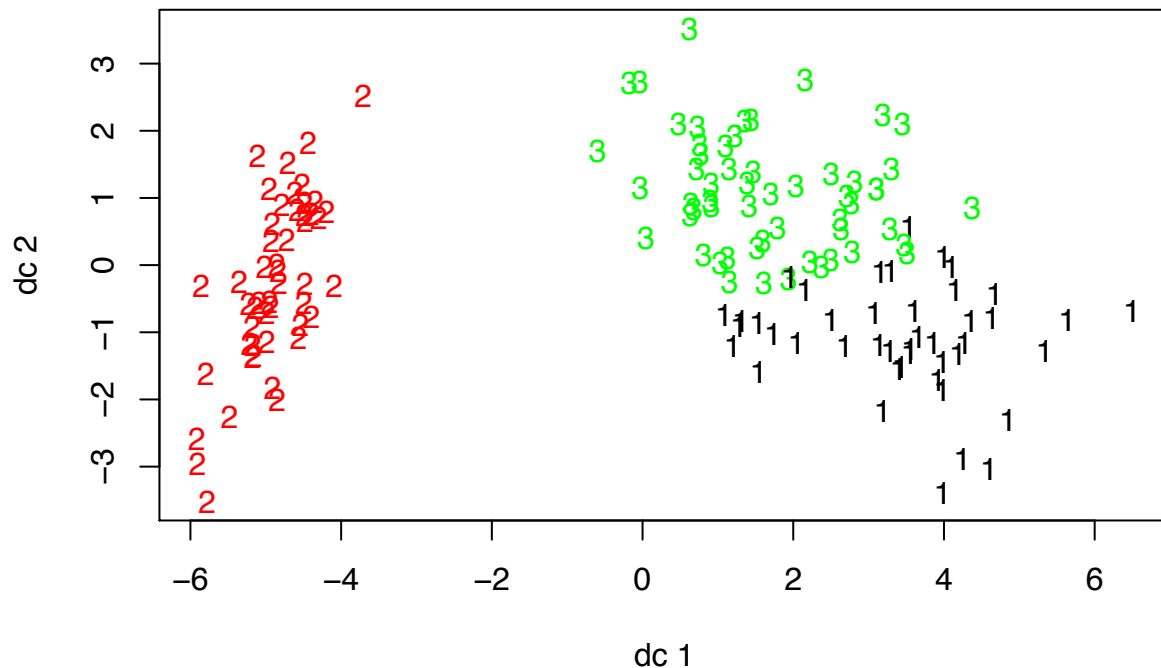
```
table3 <- table(iris$Species, kmeans_iris[,1])
ari3<- randIndex(table3); ari3
```

```
##      ARI
## 0.7302383
```

The results is 0.7302383 and it is greater than 0.5 but not a high value to say two clusterings are similar. Thus, my kmeans algorithm does not separate clusters well with unscaled iris data.

```
#scaled
scaled.iris = scale(iris[-5])
```

```
kmeans_iris.scaled <- mykmeans(scaled.iris, 3)
plotcluster(scaled.iris, as.list(kmeans_iris.scaled[,1]))
```

This time we used scaled iris data and cluster plot still does not seem to work well. There is overlapping at the boundaries of cluster group 1 and 3 and the boundary of 1 and 2 are very close. So the clusters here do not seem to be well-separated into three groups.

We can calculate the correct ratio.

```
correct_ratio4 = sum(kmeans_iris.scaled[,1] == iris[5])/length(kmeans_iris.scaled[,1])
```

So if we calculate with the above cluster results, then the correct ratio will be 0.1133333. It is less than 15%, so we can not say that it is good clustering. Also compare to the previous result with unscaled data, it increases almost two times.

We also can use the second method `randIndex`.

```
table4<- table(iris$Species, kmeans_iris.scaled[,1])
ari4<- randIndex(table4); ari4
```

```
##      ARI
## 0.5923326
```

Like previous results, ARI is 0.5923326 and it is greater than 0.5 but not a high value to say two clusterings are similar. Thus, my kmeans algorithm does not separate clusters well with scaled iris data. So the scale method does not help to cluster iris data.